

# Heart Rate Prediction Project

This notebook performs an analysis and prediction of heart rate using the Framingham dataset.

```
In [2]: # 1.Import Required Libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
import warnings
warnings.filterwarnings("ignore")
```

```
In [3]: # 2.Load Dataset
df = pd.read_csv('/Users/nikhilreddyponnala/Desktop/Skillfied Mentor/Project3/HeartRatePrediction/HeartRatePrediction.csv')
df.head()
```

```
Out[3]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp
0	1	39	4.0	0	0.0	0.0	0	0
1	0	46	2.0	0	0.0	0.0	0	0
2	1	48	1.0	1	20.0	0.0	0	0
3	0	61	3.0	1	30.0	0.0	0	0
4	0	46	3.0	1	23.0	0.0	0	0

```
In [4]: # 3.Null check
# Check for Null Values
df.isnull().sum()
```

```
Out[4]:
```

male	0
age	0
education	105
currentSmoker	0
cigsPerDay	29
BPMeds	53
prevalentStroke	0
prevalentHyp	0
diabetes	0
totChol	50
sysBP	0
diaBP	0
BMI	19
heartRate	1
glucose	388
TenYearCHD	0
dtype:	int64

```
In [5]: # 4. Fill missing values
```

```
# Fill Missing Values with Median
df.fillna(df.median(numeric_only=True), inplace=True)
```

```
In [6]: # 5. Stats & visualization

# Basic Statistics and Visualization
print(df.describe())


# Heart rate distribution
sns.histplot(df['heartRate'], kde=True)
plt.title("Distribution of Heart Rate")
plt.show()
```

	male	age	education	currentSmoker	cigsPerDay \
count	4240.000000	4240.000000	4240.000000	4240.000000	4240.000000
mean	0.429245	49.580189	1.979953	0.494104	8.944340
std	0.495027	8.572942	1.007087	0.500024	11.904777
min	0.000000	32.000000	1.000000	0.000000	0.000000
25%	0.000000	42.000000	1.000000	0.000000	0.000000
50%	0.000000	49.000000	2.000000	0.000000	0.000000
75%	1.000000	56.000000	3.000000	1.000000	20.000000
max	1.000000	70.000000	4.000000	1.000000	70.000000

	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol \
count	4240.000000	4240.000000	4240.000000	4240.000000	4240.000000
mean	0.029245	0.005896	0.310613	0.025708	236.667689
std	0.168513	0.076569	0.462799	0.158280	44.328480
min	0.000000	0.000000	0.000000	0.000000	107.000000
25%	0.000000	0.000000	0.000000	0.000000	206.000000
50%	0.000000	0.000000	0.000000	0.000000	234.000000
75%	0.000000	0.000000	1.000000	0.000000	262.000000
max	1.000000	1.000000	1.000000	1.000000	696.000000


	sysBP	diaBP	BMI	heartRate	glucose \
count	4240.000000	4240.000000	4240.000000	4240.000000	4240.000000
mean	132.354599	82.897759	25.799005	75.878774	81.600943
std	22.033300	11.910394	4.070775	12.023937	22.860340
min	83.500000	48.000000	15.540000	44.000000	40.000000
25%	117.000000	75.000000	23.077500	68.000000	72.000000
50%	128.000000	82.000000	25.400000	75.000000	78.000000
75%	144.000000	90.000000	28.032500	83.000000	85.000000
max	295.000000	142.500000	56.800000	143.000000	394.000000

	TenYearCHD
count	4240.000000
mean	0.151887
std	0.358953
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000



```
In [7]: # 6. Correlation heatmap

# Correlation Heatmap
plt.figure(figsize=(12,8))
corr = df.corr()
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```



In [8]: # 7. Feature & target definition

```
# Define Features and Target
df['heartRateClass'] = df['heartRate'].apply(lambda x: 1 if x >= 90 else 0)
X = df.drop(columns=['heartRate', 'heartRateClass'])
y = df['heartRateClass']
```

In [9]: # 8. Train/test split

```
# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

In [10]: # 9. Scaling

```
# Feature Scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

In [11]: # 10. Model training

```
# Model Training
model = LogisticRegression(class_weight='balanced')
model.fit(X_train_scaled, y_train)
```

Out[11]: **LogisticRegression**

```
LogisticRegression(class_weight='balanced')
```

In [12]: # 11. Predictions

```
# Predictions
```

```
y_pred = model.predict(X_test_scaled)
```

In [13]: # 12. Evaluation

```
# 📈 Evaluation Metrics
```

```
print(" Accuracy Score:", accuracy_score(y_test, y_pred))
```

```
print("\n\n Classification Report:\n\n", classification_report(y_test, y_pred))
```

```
print(" Confusion Matrix:\n\n", confusion_matrix(y_test, y_pred))
```

```
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, cmap='Blues', fmt='')
```

```
plt.title("Confusion Matrix")
```

```
plt.xlabel("Predicted")
```

```
plt.ylabel("Actual")
```

```
plt.show()
```

Accuracy Score: 0.6273584905660378

\n\n Classification Report:\n\n

support

precision

recall

f1-score


s

0	0.89	0.65	0.75	725
1	0.20	0.52	0.29	123

accuracy			0.63	848
macro avg	0.54	0.58	0.52	848
weighted avg	0.79	0.63	0.68	848

Confusion Matrix:\n\n [[468 257]

[ 59 64]]



## 13. Conclusion

### Conclusion

- We used the **Framingham Heart Study** dataset to predict whether an individual has a **high heart rate ( $\geq 90$  bpm)** based on various health-related features.
- After cleaning and exploring the data, we engineered a binary classification target and trained a **logistic regression model**.
- The model achieved a **reasonable accuracy**, and evaluation metrics (precision, recall, and F1-score) indicate that it can be useful for **early identification of elevated heart rate conditions**.
- To further improve model performance, we can:
  - Experiment with **other classifiers** like Random Forest, XGBoost, or SVM.
  - Perform **hyperparameter tuning** using GridSearchCV.
  - Apply **feature selection** and handle **class imbalance** more robustly.
- This study demonstrates the potential for using machine learning in **preventive cardiovascular health screening**.

Thank You

In [ ]: